

#LaUISqueQueremos



EYES ON THE TRACK

Henry Peña
2150606

Diego Medina
2150011

Motivación:

Prevenir accidentes automovilísticos causados por conductores distraídos, tales como los 85.426 accidentes con heridos y 3.406 con muertos ocurridos en Colombia en 2016. (ANSV 2019)

Enviar o leer un mensaje de texto quita sus ojos del camino por aproximadamente 5 segundos, lo suficiente como para recorrer la longitud de una cancha de fútbol al manejar a 55 millas por hora (88km/h).



*“una persona al volante en Colombia tiene 4 veces más probabilidades de morir en un accidente de tránsito que un conductor en España o Gran Bretaña”
(Revista Motor, 2019).*

Según **Fesvial, Federación Española de la Seguridad Vial**, el 45% de los accidentes se podrían prevenir si los conductores estuvieran siempre atentos a la carretera.

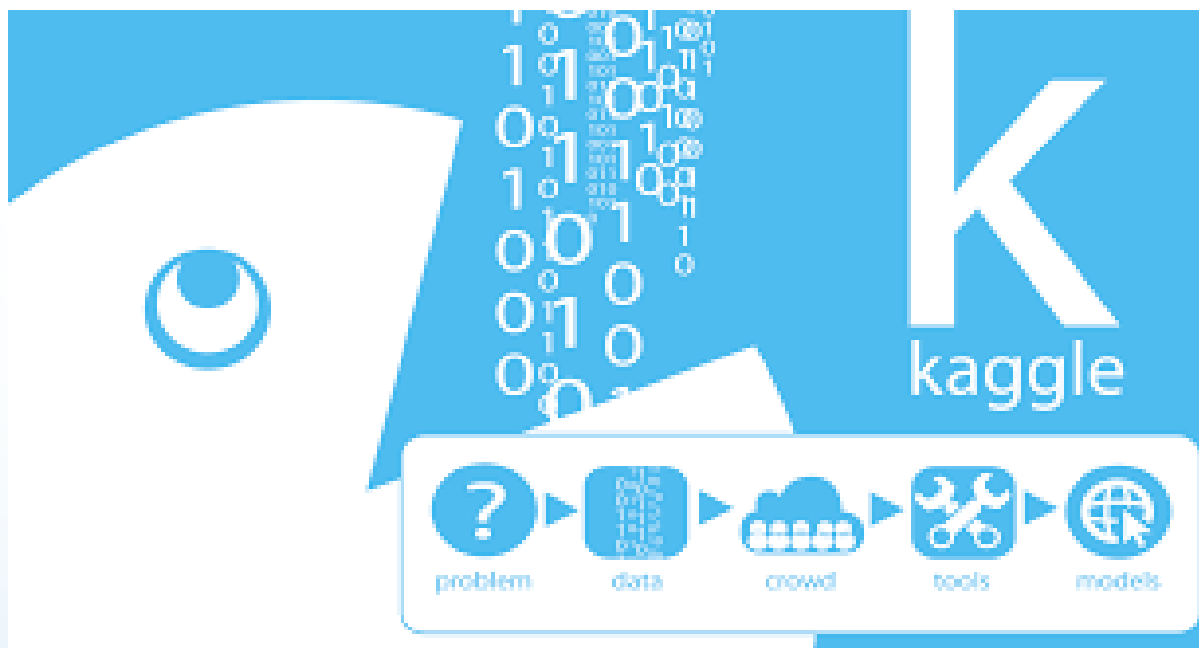
Dataset:



State Farm Distracted Driver Detection

Can computer vision spot distracted drivers?

\$65,000 · 1,440 teams · 3 years ago



El dataset que se usó para este descriptor fue obtenido de un challenge de **Kaggle** hecho por la compañía de seguros de automóviles más grande de los Estados Unidos **State Farm Insurance**.

Este dataset contiene 22400 imágenes para entrenamiento de varios conductores en un automóvil, cada cuál tiene asignada una etiqueta de la acción que están realizando (mensajes de texto, comer, hablar por teléfono, maquillaje, alcanzar detrás, etc). Hay 79.700 imágenes sin etiquetar para prueba, de distintos conductores al de entrenamiento.

Las 10 clases a predecir son:

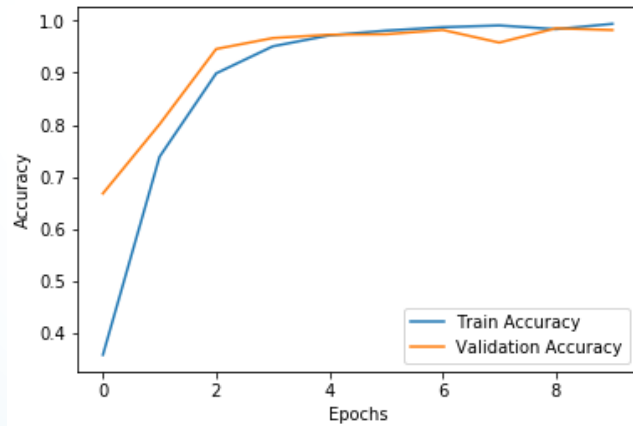
- c0: safe driving
- c1: texting - right
- c2: talking on the phone - right
- c3: texting - left
- c4: talking on the phone - left
- c5: operating the radio
- c6: drinking
- c7: reaching behind
- c8: hair and makeup
- c9: talking to passenger



DNN:

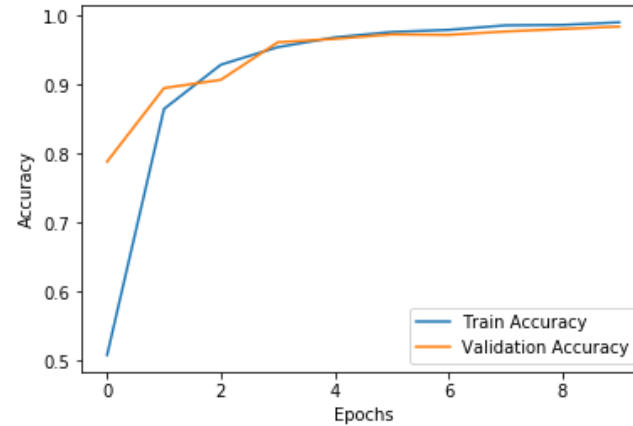
```
model.summary()
```

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 12288)	0
dense_1 (Dense)	(None, 512)	6291968
dense_2 (Dense)	(None, 256)	131328
dense_3 (Dense)	(None, 10)	2570
Total params: 6,425,866		
Trainable params: 6,425,866		
Non-trainable params: 0		




```
model5.summary()
```

Layer (type)	Output Shape	Param #
flatten_5 (Flatten)	(None, 12288)	0
dense_8 (Dense)	(None, 10)	122890
Total params: 122,890		
Trainable params: 122,890		
Non-trainable params: 0		



```
model3.summary()
```

Layer (type)	Output Shape	Param #
flatten_8 (Flatten)	(None, 12288)	0
dense_28 (Dense)	(None, 128)	1572992
dense_29 (Dense)	(None, 256)	33024
dense_30 (Dense)	(None, 512)	131584
dense_31 (Dense)	(None, 1024)	525312
Total params: 2,262,912		
Trainable params: 2,262,912		
Non-trainable params: 0		



BoW:

Data For Dictionary:

```
path_images='../datasets/driver-detections/train/'
descriptor_extractor = ORB(n_keypoints=200)
array_ORB=[]
size=100 #<1911
for i in tqdm(np.unique(data1["classname"])):
    for j in np.random.choice(data1["img"][data1.classname==i],size):
        #for j in data1["img"][data1.classname==i]:
            temp_image=cv.cvtColor(cv.imread(path_images+i+"/"+j), cv.COLOR_BGR2GRAY)
            descriptor_extractor.detect_and_extract(temp_image)
            if(len(array_ORB)==0):
                array_ORB=descriptor_extractor.descriptors
            else:
                array_ORB=np.vstack((array_ORB,descriptor_extractor.descriptors))
```

100%|██████████| 10/10 [08:17<00:00, 49.78s/it]

Se utilizaron 100 imágenes de cada clase para hacer el diccionario


Dictionary Creation:

```
: from sklearn.cluster import KMeans
v_words=70
estimator=KMeans(n_clusters=v_words).fit(array_ORB)
dictionary=estimator.cluster_centers_
print dictionary.shape
```

(70, 256)


KNN:

0.7855




Gaussian:

0.277




Random Forest:

0.703



SVC (rbf kernel):

0.7725



```

opt = keras.optimizers.SGD(lr=1e-3, momentum=0.9)
.compile(optimizer=opt, loss='sparse_categorical_crossentropy', metrics=['accuracy'])
.fit(X_train, y_train, epochs=10, batch_size=60, verbose=2, validation_data=(X_test, y_test))

```

CNN:

Model 1:


```

model_CNN = keras.models.Sequential()
model_CNN.add(keras.layers.Conv2D(20, (5, 5), activation='relu',
model_CNN.add(keras.layers.MaxPooling2D((2, 2)))
model_CNN.add(keras.layers.Conv2D(50, (5, 5), activation='relu',
model_CNN.add(keras.layers.MaxPooling2D((2, 2)))
model_CNN.add(keras.layers.Flatten())
model_CNN.add(keras.layers.Dense(500, activation='relu', kernel_i
model_CNN.add(keras.layers.Dense(10, activation='softmax'))

model_CNN.summary()

```

<matplotlib.legend.Legend at 0x29420190f08>



Model 2:


```

model_CNN2 = keras.models.Sequential()
model_CNN2.add(keras.layers.Conv2D(64, (3, 3), activation='relu',
model_CNN2.add(keras.layers.Conv2D(32, (3, 3), activation='relu',
model_CNN2.add(keras.layers.MaxPooling2D((2, 2)))
model_CNN2.add(keras.layers.Flatten())
model_CNN2.add(keras.layers.Dense(512, activation='relu', kernel
model_CNN2.add(keras.layers.Dense(256, activation='relu', kernel
model_CNN2.add(keras.layers.Dense(128, activation='relu', kernel
model_CNN2.add(keras.layers.Dense(10, activation='softmax'))

model_CNN2.summary()

```

<matplotlib.legend.Legend at 0x29420644d88>



Model 3:

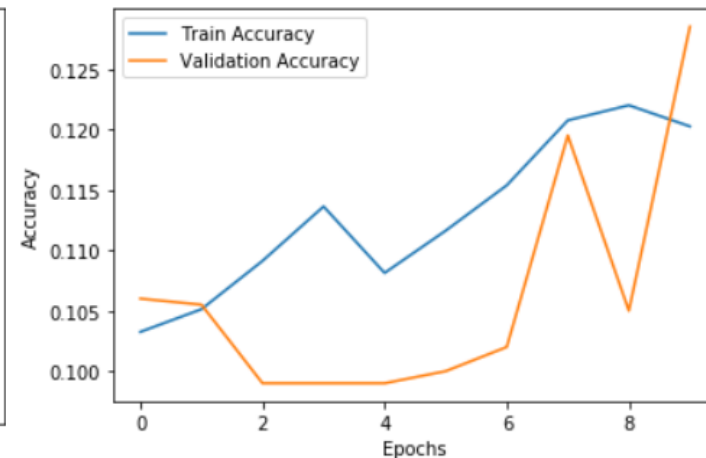
```

model_CNN3 = keras.models.Sequential()
model_CNN3.add(keras.layers.Conv2D(32, (3, 3), ke
model_CNN3.add(keras.layers.MaxPooling2D((2, 2)))
model_CNN3.add(keras.layers.Dropout(0.5))
model_CNN3.add(keras.layers.Conv2D(64, (3, 3), ke
model_CNN3.add(keras.layers.MaxPooling2D((2, 2)))
model_CNN3.add(keras.layers.Dropout(0.5))
model_CNN3.add(keras.layers.Conv2D(128, (3, 3), k
model_CNN3.add(keras.layers.MaxPooling2D((8, 8)))
model_CNN3.add(keras.layers.Dropout(0.5))
model_CNN3.add(keras.layers.Flatten())
model_CNN3.add(keras.layers.Dense(10, activation=

model_CNN3.summary()

```

<matplotlib.legend.Legend at 0x2943c088e48>



Modelo Escogido:

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 60, 60, 20)	1520
max_pooling2d (MaxPooling2D)	(None, 30, 30, 20)	0
conv2d_1 (Conv2D)	(None, 26, 26, 50)	25050
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 50)	0
flatten (Flatten)	(None, 8450)	0
dense (Dense)	(None, 500)	4225500
dense_1 (Dense)	(None, 10)	5010

Total params: 4,257,080

Trainable params: 4,257,080

Non-trainable params: 0

Entrenamiento:

```
[ ] history = model_CNN.fit(X_train,y_train, epochs=10, batch_size=160, verbose=2,validation_data=(X_test, y_test))
```



Train on 17939 samples, validate on 4485 samples

Epoch 1/10

17939/17939 - 6s - loss: 1.9033 - acc: 0.3476 - val_loss: 1.5848 - val_acc: 0.4531

Epoch 2/10

17939/17939 - 2s - loss: 1.1783 - acc: 0.6422 - val_loss: 0.8958 - val_acc: 0.7358

Epoch 3/10

17939/17939 - 2s - loss: 0.6903 - acc: 0.8123 - val_loss: 0.6564 - val_acc: 0.8004

Epoch 4/10

17939/17939 - 2s - loss: 0.4387 - acc: 0.8922 - val_loss: 0.3983 - val_acc: 0.8965

Epoch 5/10

17939/17939 - 2s - loss: 0.2887 - acc: 0.9334 - val_loss: 0.2954 - val_acc: 0.9240

Epoch 6/10

17939/17939 - 2s - loss: 0.2183 - acc: 0.9513 - val_loss: 0.2125 - val_acc: 0.9494

Epoch 7/10

17939/17939 - 2s - loss: 0.1627 - acc: 0.9648 - val_loss: 0.1710 - val_acc: 0.9610

Epoch 8/10

17939/17939 - 2s - loss: 0.1300 - acc: 0.9722 - val_loss: 0.1440 - val_acc: 0.9648

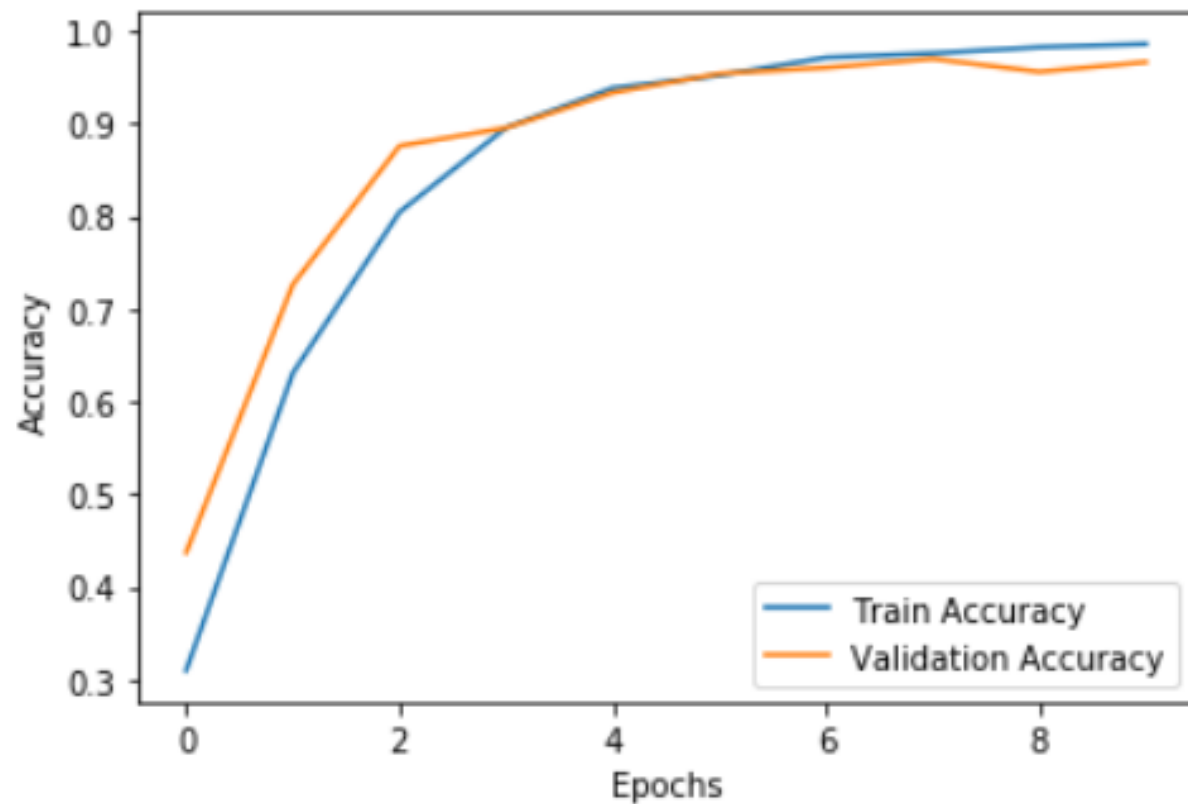
Epoch 9/10

17939/17939 - 2s - loss: 0.1035 - acc: 0.9784 - val_loss: 0.1173 - val_acc: 0.9715


Epoch 10/10

17939/17939 - 2s - loss: 0.0842 - acc: 0.9823 - val_loss: 0.1127 - val_acc: 0.9710

Accuracy/epoch:



8.000 images Train
2.000 images Test

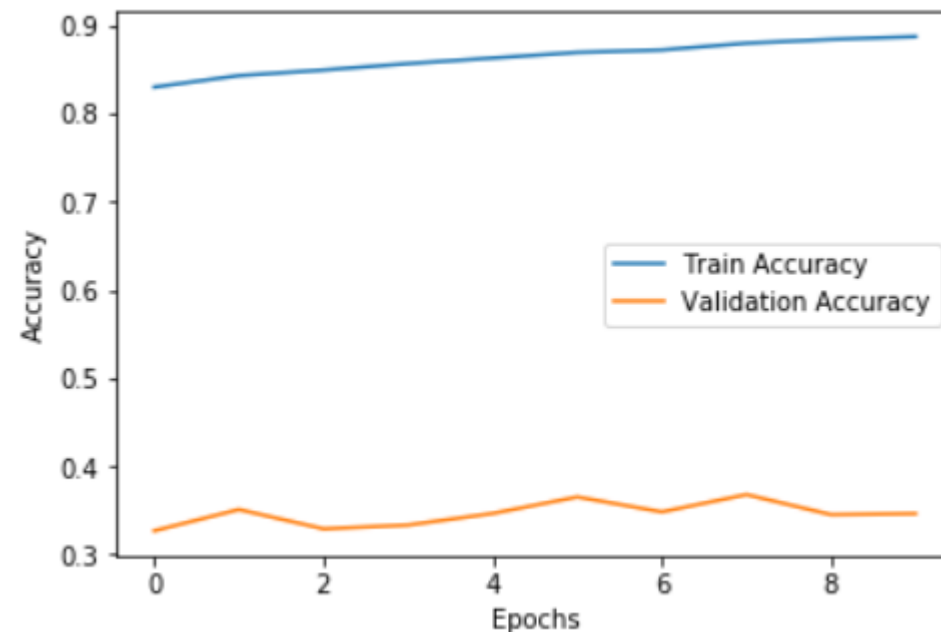



17.939 images Train
4.484 images Test

Transfer Learning (Fully Connected train):

MobileNet:

Vgg19:



Algunos ejemplos:

texting - left



reaching behind



drinking



talking on the phone - right



talking to passenger



Safe Driving



Kaggle Submission:

```
[ ] predicts=[]
```

```
[ ] for i in tqdm(range(0,19974)):
    temp_img=cv.resize(cv.cvtColor(cv.imread(path_test_images+'/' +
    temp_predict=model_CNN.predict(np.array([temp_img])))
    predicts.append(np.append(test_images_names[i],temp_predict))
```

```
100%|██████████| 19974/19974 [1:19:43<00:00, 4.06it/s]
```

```
[ ] for i in tqdm(range(19974, 39948)):
    temp_img=cv.resize(cv.cvtColor(cv.imread(path_test_images+'/' +
    temp_predict=model_CNN.predict(np.array([temp_img])))
    predicts.append(np.append(test_images_names[i],temp_predict))
```

```
100%|██████████| 19974/19974 [1:17:49<00:00, 4.03it/s]
```

```
[ ] for i in tqdm(range(39948, 59922)):
    temp_img=cv.resize(cv.cvtColor(cv.imread(path_test_images+'/' +
    temp_predict=model_CNN.predict(np.array([temp_img])))
    predicts.append(np.append(test_images_names[i],temp_predict))
```

```
100%|██████████| 19974/19974 [1:12:52<00:00, 4.10it/s]
```

```
[ ] for i in tqdm(range(59922, 79896)):
    temp_img=cv.resize(cv.cvtColor(cv.imread(path_test_images+'/' +
    temp_predict=model_CNN.predict(np.array([temp_img])))
    predicts.append(np.append(test_images_names[i],temp_predict))
```

```
100%|██████████| 19974/19974 [1:11:08<00:00, 4.78it/s]
```

Wait time
0 seconds

Execution time
1 seconds

Score
17.31954



Universidad
Industrial de
Santander

www.uis.edu.co

**Repositorio del
proyecto:**

<https://github.com/pecons/proyectoCV>



Gracias
por su
atención